LINEAR LENSES IN HASKELL

I always believed that the main problems in designing a programming language were resource management and concurrency–and the two are related. If you can track ownership of resources, you can be sure that no synchronization is needed when there's a single owner.

I've been evangelizing resource management for a long time, first in C++, and then in D (see Appendix 3). I was happy to see it implemented in Rust as ownership types, and I'm happy to see it coming to Haskell as linear types.

Haskell has essentially solved the concurrency and parallelism problems by channeling mutation to dedicated monads, but resource management has always been part of the awkward squad. The main advantage of linear types in Haskell, other than dealing with external resources, is to allow safe mutation and non-GC memory management. This could potentially lead to substantial performance gains.

This post is based on very informative discussions I had with Arnaud Spiwack, who explained to me the work he'd done on linear types and linear lenses, some of it never before documented.

The PDF version of this post together with complete Haskell code is available on GitHub.

1. Linear types

What is a linear function? The short answer is that a linear function $a \multimap b$ "consumes" its argument *exactly once*. This is not the whole truth, though, because we also have linear identity $id_a\colon a \multimap a$, which ostensibly does not consume a. The long answer is that a linear function consumes its argument exactly once if it itself is consumed exactly once, and its result is consumed exactly once.

What remains to define is what it means to be consumed. A function is consumed when it's applied to its argument. A base value like `Int` or `Char` is consumed when it's evaluated, and an algebraic data type is consumed when it's pattern-matched and all its matched components are consumed.

For instance, to consume a linear pair `(a, b)`, you pattern-match it and then consume both `a` and `b`. To consume `Either a b`, you pattern-match it and consume the matched component, either `a` or `b`, depending on which branch was taken.

As you can see, except for the base values, a linear argument is like a hot potato: you "consume" it by passing it to somebody else.

So where does the buck stop? This is where the magic happens: Every resource comes with a special primitive that gets rid of it. A file handle gets

closed, memory gets deallocated, an array gets frozen, and Frodo throws the ring into the fires of Mount Doom.

To notify the type system that the resource has been destroyed, a linear function will return a value inside the special *unrestricted* type `Ur`. When this type is pattern-matched, the original resource is finally destroyed.

For instance, for linear arrays, one such primitive is `toList`:

$$toList \colon \mathrm{Array}\ a \multimap \mathrm{Ur}\,[a]$$

In Haskell, we annotate the linear arrows with multiplicity 1:

```
toList :: Array a %1-> Ur  [a]
```

Similarly, magic is used to create the resource in the first place. For arrays, this happens inside the primitive `fromList`.

$$fromList \colon [a] \to (\mathrm{Array}\ a \multimap \mathrm{Ur}\ b) \multimap \mathrm{Ur}\ b$$

or using Haskell syntax:

```
fromList :: [a] -> (Array a %1-> Ur b) %1-> Ur b
```

The kind of resource management I advertised in C++ was scope based. A resource was encapsulated in a smart pointer that was automatically destroyed at scope exit.

With linear types, the role of the scope is played by a user-provided linear function; here, the continuation:

```
(Array a %1 -> Ur b)
```

The primitive `fromList` promises to consume this user-provided function exactly once and to return its unrestricted result. The client is obliged to consume the array exactly once (e.g., by calling `toList`). This obligation is encoded in the type of the continuation accepted by `fromList`.

2. Linear lens: The existential form

A lens abstracts the idea of focusing on a part of a larger data structure. It is used to access or modify its *focus*. An existential form of a lens consists of two functions: one splitting the source into the focus and the residue; and the other replacing the focus with a new value, and creating a new whole. We don't care about the actual type of the residue so we keep it as an existential.

The way to think about a linear lens is to consider its source as a resource. The act of splitting it into a focus and a residue is destructive: it consumes its source to produce two new resources. It splits one hot potato `s` into two hot potatoes: the residue `c` and the focus `a`.

Conversely, the part that rebuilds the target `t` must consume both the residue `c` and the new focus `b`.

We end up with the following Haskell implementation:

```haskell
data LinLensEx a b s t where
  LinLensEx :: (s %1-> (c, a)) ->
               ((c, b) %1-> t) -> LinLensEx a b s t
```

A Haskell existential type corresponds to a categorical coend, so the above definition is equivalent to:

$$Labst = \int^c (s \multimap c \otimes a) \times (c \otimes b \multimap t)$$

I use the lollipop notation for the hom-set in a monoidal category with a tensor product $\otimes$.

The important property of a monoidal category is that its tensor product doesn't come with a pair of projections; and the unit object is not terminal. In particular, a morphism $s \multimap c \otimes a$ cannot be decomposed into a product of two morphisms $(s \multimap c) \times (s \multimap a)$.

However, in a *closed* monoidal category we can curry a mapping out of a tensor product:

$$c \otimes b \multimap t \cong c \multimap (b \multimap t)$$

We can therefore rewrite the existential lens as:

$$Labst \cong \int^c (s \multimap c \otimes a) \times (c \multimap (b \multimap t))$$

and then apply the co-Yoneda lemma to get:

$$s \multimap \big((b \multimap t) \otimes a\big)$$

Unlike the case of a standard lens, this form cannot be separated into a get/set pair.

The intuition is that a linear lens lets you consume the object s, but it leaves you with the obligation to consume both the setter $b \multimap t$ and the focus a. You can't just extract a alone, because that would leave a gaping hole in your object. You have to plug it in with a new object b, and that's what the setter lets you do.

Here's the Haskell translation of this formula (conventionally, with the pairs of arguments reversed):

```haskell
type LinLens s t a b = s %1-> (b %1-> t, a)
```

The Yoneda shenanigans translate into a pair of Haskell functions. Notice that, just like in the co-Yoneda trick, the existential c is replaced by the linear function $b \multimap t$.

```haskell
fromLinLens :: forall s t a b.
  LinLens s t a b -> LinLensEx a b s t
fromLinLens h = LinLensEx f g
  where
    f :: s %1-> (b %1-> t, a)
```

```
    f = h
    g :: (b %1-> t, b) %1-> t
    g (set, b) = set b
```

The inverse mapping is:

```
toLinLens :: LinLensEx a b s t -> LinLens s t a b
toLinLens (LinLensEx f g) s =
  case f s of
    (c, a) -> (\b -> g (c, b), a)
```

3. PROFUNCTOR REPRESENTATION

Every optic comes with a profunctor representation and the linear lens is no exception. Categorically speaking, a profunctor is a functor from the product category $\mathcal{C}^{op} \times \mathcal{C}$ to **Set**. It maps pairs of object to sets, and pairs of morphisms to functions. Since we are in a monoidal category, the morphisms are linear functions, but the mappings between sets are regular functions (see Appendix 1). Thus the action of the profunctor p on morphisms is a function:

$$(a' \multimap a) \to (b \multimap b') \to pab \to pa'b'$$

In Haskell:

```
class Profunctor p where
  dimap :: (a' %1-> a) -> (b %1-> b') -> p a b -> p a' b'
```

A Tambara module (a.k.a., a strong profunctor) is a profunctor equipped with the following mapping:

$$\alpha_{abc} \colon pab \to p(c \otimes a)(c \otimes b)$$

natural in a and b, dintural in c. In Haskell, this translates to a polymorphic function:

```
class (Profunctor p) => Tambara p where
  alpha :: forall a b c. p a b -> p (c, a) (c, b)
```

The linear lens $Labst$ is itself a Tambara module, for fixed ab. To show this, let's construct a mapping:

$$\alpha_{stc} \colon Labst \to Lab(c \otimes s)(c \otimes t)$$

Expanding the definition:

$$\int^{c''} (s \multimap c'' \otimes a) \times (c'' \otimes b \multimap t) \to \int^{c'} (c \otimes s \multimap c' \otimes a) \times (c' \otimes b \multimap c \otimes t)$$

By cocontinuity of the hom-set in **Set**, a mapping out of a coend is equivalent to an end:

$$\int_{c''} \left((s \multimap c'' \otimes a) \times (c'' \otimes b \multimap t) \to \int^{c'} (c \otimes s \multimap c' \otimes a) \times (c' \otimes b \multimap c \otimes t)\right)$$

Given a pair of linear arrows on the left we want to construct a coend on the right. We can do it by first lifting both arrow by $(c \otimes -)$. We get:

$$(c \otimes s \multimap c \otimes c'' \otimes a) \times (c \otimes c'' \otimes b \multimap c \otimes t)$$

We can inject them into the coend on the right at $c' = c \otimes c''$.

In Haskell, we construct the instance of the `Profunctor` class for the linear lens:

```
instance Profunctor (LinLensEx a b) where
  dimap f' g' (LinLensEx f g) = LinLensEx (f . f') (g' . g)
```

and the instance of `Tambara`:

```
instance Tambara (LinLensEx a b) where
  alpha (LinLensEx f g) =
    LinLensEx (unassoc . second f) (second g . assoc)
```

Linear lenses can be composed and there is an identity linear lens:

$$id_{ab} \colon Labab = \int^c (a \multimap c \otimes a) \times (c \otimes b \multimap b)$$

given by injecting the pair (id_a, id_b) at $c = I$, the monoidal unit.

In Haskell, we can construct the identity lens using the left unitor (see Appendix 1):

```
idLens :: LinLensEx a b a b
idLens = LinLensEx unlunit lunit
```

The profunctor representation of a linear lens is given by an end over Tambara modules:

$$Labst \cong \int_{p:Tamb} pab \to pst$$

In Haskell, this translates to a type of functions polymorphic in Tambara modules:

```
type PLens a b s t = forall p. Tambara p => p a b -> p s t
```

The advantage of this representation is that it lets us compose linear lenses using simple function composition.

Here's the categorical proof of the equivalence. Left to right: Given a triple: $(c, f \colon s \multimap c \otimes a, g \colon c \otimes b \multimap t)$, we construct:

$$pab \xrightarrow{\alpha_{abc}} p(c \otimes a)(c \otimes b) \xrightarrow{pfg} pst$$

Conversely, given a polymorphic (in Tambara modules) function $pab \to pst$, we can apply it to the identity optic id_{ab} and obtain $Labst$.

In Haskell, this equivalence is witnessed by the following pair of functions:

```haskell
fromPLens :: PLens a b s t -> LinLensEx a b s t
fromPLens f = f idLens
```

```haskell
toPLens :: LinLensEx a b s t -> PLens a b s t
toPLens (LinLensEx f g) pab = dimap f g (alpha pab)
```

4. van Laarhoven representation

Similar to regular lenses, linear lenses have a functor-polymorphic van Laarhoven encoding. The difference is that we have to use endofunctors in the monoidal subcategory, where all arrows are linear:

```haskell
class Functor f where
  fmap :: (a %1-> b) %1-> f a %1-> f b
```

Just like regular Haskell functors, linear functors are strong. This follows from the fact that a closed monoidal category is self-enriched. We define strength as:

```haskell
strength :: Functor f => (a, f b) %1-> f (a, b)
strength (a, fb) = fmap (eta a) fb
```

where `eta` is the unit of the currying adjunction (see Appendix 1).

With this definition, the van Laarhoven encoding of linear lenses is:

```haskell
type VLL s t a b = forall f. Functor f =>
    (a %1-> f b) -> (s %1-> f t)
```

The equivalence of the two encodings is witnessed by a pair of functions:

```haskell
toVLL :: LinLens s t a b -> VLL s t a b
toVLL lns f = fmap apply . strength . second f . lns
```

```haskell
fromVLL :: forall s t a b. VLL s t a b -> LinLens s t a b
fromVLL vll s = unF (vll (F id) s)
```

Here, the functor `F` is defined as a linear pair (a tensor product):

```haskell
data F a b x where
    F :: (b %1-> x) %1-> a %1-> F a b x
unF :: F a b x %1-> (b %1-> x, a)
unF (F bx a) = (bx, a)
```

with the obvious implementation of `fmap`

```
instance Functor (F a b) where
  fmap f (F bx a) = F (f . bx) a
```

You can find the categorical derivation of van Laarhoven representation in Appendix 2.

5. Linear optics

Linear lenses are but one example of more general linear optics. Linear optics are defined by the action of a monoidal category $\mathcal{M}$ on (possibly the same) monoidal category $\mathcal{C}$:

$$\bullet\colon \mathcal{M} \times \mathcal{C} \to \mathcal{C}$$

In particular, one can define linear prisms and linear traversals using actions by a coproduct or a power series.

The existential form is given by:

$$Oabst = \int^{m\colon \mathcal{M}} (s \multimap m \bullet a) \times (m \bullet b \multimap t)$$

There is a corresponding Tambara representation, with the following Tambara structure:

$$\alpha_{abm}\colon pab \to p(m \bullet a)(m \bullet b)$$

Incidentally, the two hom-sets in the definition of the optic can come from two different categories, so it's possible to mix linear and non-linear arrows in one optic.

6. Appendix: 1 Closed monoidal category in Haskell

With the advent of linear types we now have two main categories lurking inside Haskell. They have the same objects–Haskell types– but the monoidal category has fewer arrows. These are the linear arrows $a \multimap b$. They can be composed:

```
(.) :: (b %1-> c) %1-> (a %1-> b) %1-> a %1 -> c
(f . g) x = f (g x)
```

and there is an identity arrow for every object:

```
id :: a %1-> a
id a = a
```

In general, a tensor product in a monoidal category is a bifunctor: $\mathcal{C} \times \mathcal{C} \to \mathcal{C}$. In Haskell, we identify the tensor product $\otimes$ with the built-in product (a, b). The difference is that, within the monoidal category, this product doesn't have projections. There is no linear arrow $(a, b) \multimap a$ or $(a, b) \multimap b$. Consequently, there is no diagonal map $a \multimap (a, a)$, and the unit object $()$ is not terminal: there is no arrow $a \multimap ()$.

We define the action of a bifunctor on a pair of linear arrows entirely within the monoidal category:

```haskell
class Bifunctor p where
    bimap :: (a %1-> a') %1-> (b %1-> b') %1->
              p a b %1-> p a' b'
    first :: (a %1-> a') %1-> p a b %1-> p a' b
    first f = bimap f id
    second :: (b %1-> b') %1-> p a b %1-> p a b'
    second = bimap id
```

The product is itself an instance of this linear bifunctor:

```haskell
instance Bifunctor (,) where
    bimap f g (a, b) = (f a, g b)
```

The tensor product has to satisfy coherence conditions–associativity and unit laws:

```haskell
assoc :: ((a, b), c) %1-> (a, (b, c))
assoc ((a, b), c) = (a, (b, c))
unassoc :: (a, (b, c)) %1-> ((a, b), c)
unassoc (a, (b, c)) = ((a, b), c)
```

```haskell
lunit :: ((), a) %1-> a
lunit ((), a) = a
unlunit :: a %1-> ((), a)
unlunit a = ((), a)
```

In Haskell, the type of arrows between any two objects is also an object. A category in which this is true is called closed. This identification is the consequence of the currying adjunction between the product and the function type. In a closed monoidal category, there is a corresponding adjunction between the tensor product and the object of linear arrows. The mapping out of a tensor product is equivalent to the mapping into the function object. In Haskell, this is witnessed by a pair of mappings:

```haskell
curry :: ((a, b) %1-> c) %1-> (a %1-> (b %1-> c))
curry f x y = f (x, y)

uncurry :: (a %1-> (b %1-> c)) %1-> ((a, b) %1-> c)
uncurry f (x, y) = f x y
```

Every adjunction also defines a pair of unit and counit natural transformations:

```haskell
eta :: a %1-> b %1-> (a, b)
eta a b = (a, b)
```

```
apply :: (a %1-> b, a) %-> b
apply (f, a) = f a
```

We can, for instance, use the unit to implement a point-free mapping of lenses:

```
toLinLens :: LinLensEx a b s t -> LinLens s t a b
toLinLens (LinLensEx f g) = first ((g .) . eta) . f
```

Finally, a note about the Haskell definition of a profunctor:

```
class Profunctor p where
  dimap :: (a' %1-> a) -> (b %1-> b') -> p a b -> p a' b'
```

Notice the mixing of two types of arrows. This is because a profunctor is defined as a mapping $\mathcal{C}^{op} \times \mathcal{C} \to \mathbf{Set}$. Here, $\mathcal{C}$ is the monoidal category, so the arrows in it are linear. But `p a b` is just a set, and the mapping `p a b -> p a' b'` is just a regular function. Similarly, the type:

```
(a' %1-> a)
```

is not treated as an object in $\mathcal{C}$ but rather as a set of linear arrows. In fact this hom-set is itself a profunctor:

```
newtype Hom a b = Hom (a %1-> b)

instance Profunctor Hom where
  dimap f g (Hom h) = Hom (g . h . f)
```

As you might have noticed, there are many definitions that extend the usual Haskel concepts to linear types. Since it makes no sense to re-introduce, and give new names to each of them, the linear extensions are written using multiplicity polymorphism. For instance, the most general currying function is written as:

```
curry :: ((a, b) %p -> c) %q -> a %p -> b %p -> c
```

covering four different combinations of multiplicities.

7. Appendix 2: van Laarhoven representation

We start by defining functorial strength in a monoidal category:

$$\sigma_{ab}\colon a \otimes Fb \multimap F(a \otimes b)$$

In a closed monoidal category every functor is strong. To begin with, we can curry σ. Thus we have to construct:

$$a \multimap (Fb \multimap F(a \otimes b))$$

We have at our disposal the counit of the currying adjunction:

$$\eta_{ab} \colon a \multimap (b \multimap a \otimes b)$$

We can apply η_{ab} to a and lift the resulting map $(b \multimap a \otimes b)$ to arrive at $Fb \multimap F(a \otimes b)$.

Now let's write the van Laarhoven representation as the end of the mapping of two linear hom-sets:

$$\int_{F\colon [\mathcal{C},\mathcal{C}]} (a \multimap Fb) \to (s \multimap Ft)$$

We use the Yoneda lemma to replace $a \multimap Fb$ with a set of natural transformations, written as an end over x:

$$\int_F \int_x \left((b \multimap x) \multimap (a \multimap Fx) \right) \to (s \multimap Ft)$$

We can uncurry it:

$$\int_F \int_x \left((b \multimap x) \otimes a \multimap Fx \right) \to (s \multimap Ft)$$

and apply the ninja-Yoneda lemma in the functor category to get:

$$s \multimap \left((b \multimap t) \otimes a \right)$$

Here, the ninja-Yoneda lemma operates on higher-order functors, such as $\Phi_{st} F = (s \multimap Ft)$. It can be written as:

$$\int_F \int_x (Gx \multimap Fx) \to \Phi_{st} F \cong \Phi_{st} G$$

8. Appendix 3: My resource management curriculum

These are some of my blog posts and articles about resource management and its application to concurrent programming.

(1) Strong Pointers and Resource Management in C++,
- Part 1, 1999
- Part 2, 2000
(2) Walking Down Memory Lane, 2005 (with Andrei Alexandrescu)
(3) unique ptr–How Unique is it?, 2009
(4) Unique Objects, 2009
(5) Race-free Multithreading, 2009
- Part 1: Ownership
- Part 2: Owner Polymorphism
(6) Edward C++ hands, 2013